

microbiology

April 26, 2025

```
[37]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import timedelta
import gc # Garbage Collector

# Import GCS library
from google.cloud import storage

from sklearn.model_selection import train_test_split, cross_val_score, \
    ↳StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
# Updated imports for multiclass evaluation
from sklearn.metrics import classification_report, confusion_matrix, \
    ↳accuracy_score, f1_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder # \
    ↳Added LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from collections import Counter
```

```
[39]: # Define time windows and lookback periods
TIME_WINDOW_HOURS = 24
PRIOR_ANTIBIOTIC_LOOKBACK_DAYS = 30
TOP_N_PAIRS = 10
N_SPLITS_CV = 5
```

```
[2]: # Define required CSV files and their paths
gcs_paths = {
    "micro": "gs://msai/aih/MICROBIOLOGYEVENTS.csv.gz",
    "patients": "gs://msai/aih/PATIENTS.csv.gz",
    "admissions": "gs://msai/aih/ADMISSIONS.csv.gz",
    "labevents": "gs://msai/aih/LABEVENTS.csv.gz",
    "chartevents": "gs://msai/aih/CHARTEVENTS.csv.gz",
```

```

    "prescriptions": "gs://msai/aih/PRESCRIPTIONS.csv.gz",
    "diagnoses": "gs://msai/aih/DIAGNOSES_ICD.csv.gz"
}

```

```

[3]: # Define dtypes for memory efficiency
# Using appropriate integer types and 'category' can save significant memory
dtypes_micro = {
    'SUBJECT_ID': 'int32', 'HADM_ID': 'float32',
    'SPEC_ITEMID': 'float32', 'ORG_ITEMID': 'float32', 'AB_ITEMID': 'float32',
    'ISOLATE_NUM': 'float32', 'DILUTION_VALUE': 'float32',
    'SPEC_TYPE_DESC': 'category', 'ORG_NAME': 'category', 'AB_NAME': 'category',
    'INTERPRETATION': 'category'
}
dtypes_patients = {'SUBJECT_ID': 'int32', 'GENDER': 'category'}
dtypes_admissions = {
    'SUBJECT_ID': 'int32', 'HADM_ID': 'int32',
    'ADMISSION_TYPE': 'category', 'ETHNICITY': 'category'
}
dtypes_labevents = {
    'SUBJECT_ID': 'int32', 'HADM_ID': 'float32', 'ITEMID': 'int32',
    'VALUE': 'object', 'VALUENUM': 'float32', 'VALUEUOM': 'category', 'FLAG':
    ↪ 'category'
}
dtypes_chartevents = {
    'SUBJECT_ID': 'int32', 'HADM_ID': 'float32', 'ICUSTAY_ID': 'float32',
    ↪ 'ITEMID': 'int32',
    'VALUE': 'object', 'VALUENUM': 'float32', 'VALUEUOM': 'category', 'WARNING':
    ↪ 'float32',
    'ERROR': 'float32', 'RESULTSTATUS': 'category', 'STOPPED': 'category'
}
dtypes_prescriptions = {
    'SUBJECT_ID': 'int32', 'HADM_ID': 'int32', 'ICUSTAY_ID': 'float32',
    'DRUG_TYPE': 'category', 'DRUG': 'category', 'ROUTE': 'category'
}
dtypes_diagnoses = {
    'SUBJECT_ID': 'int32', 'HADM_ID': 'int32', 'SEQ_NUM': 'float32',
    ↪ 'ICD9_CODE': 'category'
}

```

```

[40]: def load_csv_from_gcs(gcs_path, dtypes=None, parse_dates=None, usecols=None):
    df = pd.read_csv(gcs_path, dtype=dtypes, parse_dates=parse_dates,
    ↪ usecols=usecols, compression='gzip', low_memory=False)
    print(f"Loaded {len(df)} rows.")
    return df

```

```

[41]: df_micro = load_csv_from_gcs(gcs_paths['micro'], dtypes=dtypes_micro,
    ↪ parse_dates=['CHARTTIME', 'CHARTDATE'])

```

Loaded 631726 rows.

```
[42]: # Common values are R, S, I (Intermediate).
RELEVANT_INTERPRETATIONS = ['R', 'S', 'I']
```

```
[43]: # Filter for relevant interpretations and non-null org/ab names
df_micro_filtered = df_micro[
    df_micro['INTERPRETATION'].isin(RELEVANT_INTERPRETATIONS) & # Use the list
    df_micro['ORG_NAME'].notna() &
    df_micro['AB_NAME'].notna()
].copy()
```

```
[44]: # Count pairs based on these interpretations
top_pairs_df = df_micro_filtered.groupby(['ORG_NAME', 'AB_NAME']).size().
    ↪reset_index(name='PAIR_COUNT')
top_pairs_df = top_pairs_df.sort_values('PAIR_COUNT', ascending=False).
    ↪head(TOP_N_PAIRS)
```

/var/tmp/ipykernel_16006/4283320187.py:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
top_pairs_df = df_micro_filtered.groupby(['ORG_NAME',
'AB_NAME']).size().reset_index(name='PAIR_COUNT')
```

```
[45]: top_pairs_df.head(10)
```

```
[45]:
```

	ORG_NAME	AB_NAME	PAIR_COUNT
9620	STAPH AUREUS COAG +	OXACILLIN	8546
9614	STAPH AUREUS COAG +	GENTAMICIN	7555
9616	STAPH AUREUS COAG +	LEVOFLOXACIN	7523
9613	STAPH AUREUS COAG +	ERYTHROMYCIN	7247
9611	STAPH AUREUS COAG +	CLINDAMYCIN	5659
9626	STAPH AUREUS COAG +	TETRACYCLINE	4699
9621	STAPH AUREUS COAG +	PENICILLIN	4620
9629	STAPH AUREUS COAG +	VANCOMYCIN	4614
9625	STAPH AUREUS COAG +	RIFAMPIN	4434
4064	ESCHERICHIA COLI	GENTAMICIN	4220

```
[12]: # Load tables that don't depend on the specific pair first to avoid reloading
    ↪in loop
print("\nLoading base static tables...")
df_patients = load_csv_from_gcs(gcs_paths['patients'], dtypes=dtypes_patients,
    ↪parse_dates=['DOB'])
df_admissions = load_csv_from_gcs(gcs_paths['admissions'],
    ↪dtypes=dtypes_admissions, parse_dates=['ADMITTIME', 'DISCHTIME'])
df_diagnoses_full = load_csv_from_gcs(gcs_paths['diagnoses'],
    ↪dtypes=dtypes_diagnoses)
```

```
df_prescriptions_full = load_csv_from_gcs(gcs_paths['prescriptions'],
↳ dtypes=dtypes_prescriptions, parse_dates=['STARTDATE', 'ENDDATE'])
```

Loading base static tables...
 Loaded 46520 rows.
 Loaded 58976 rows.
 Loaded 651047 rows.
 Loaded 4156450 rows.

```
[13]: # --- Load Large Event Tables Once (very memory intensive) ---
# loading these within the loop because memory is constrained.
print("\nLoading large event tables (may take time and memory)...")
df_labevents_full = load_csv_from_gcs(gcs_paths['labevents'],
↳ dtypes=dtypes_labevents, parse_dates=['CHARTTIME'])
# df_chartevents_full = load_csv_from_gcs(gcs_paths['chartevents'],
↳ dtypes=dtypes_chartevents, parse_dates=['CHARTTIME'])
```

Loading large event tables (may take time and memory)...
 Loaded 27854055 rows.

```
[15]: # Check if all dataframes loaded successfully
if any(df is None for df in [df_patients, df_admissions, df_diagnoses_full,
↳ df_prescriptions_full, df_labevents_full]):
    print("One or more required base CSV files failed to load.")
```

```
[18]: # --- Initialize Label Encoder for the Target Variable ---
target_encoder = LabelEncoder()
target_encoder.fit(RELEVANT_INTERPRETATIONS)
print(f"Target classes encoded as: {dict(zip(target_encoder.classes_,
↳ target_encoder.transform(target_encoder.classes_)))}")
```

Target classes encoded as: {np.str_('I'): np.int64(0), np.str_('R'): np.int64(1), np.str_('S'): np.int64(2)}

```
[51]: results_store = dict()
```

```
[52]: # --- Loop Through Top Pairs ---
print("\n--- Starting Analysis Loop for Top Pairs ---")
for index, row in top_pairs_df.iterrows():
    TARGET_ORG = row['ORG_NAME']
    TARGET_AB_NAME = row['AB_NAME']
    pair_label = f"{TARGET_ORG} / {TARGET_AB_NAME}"
    print(f"\n===== Processing Pair {index+1}/{TOP_N_PAIRS}:
↳ {pair_label} =====")

    # --- 1b. Data Processing for the current pair ---
```

```

print(f"\nProcessing data for {pair_label}...")

# 1. Filter Microbiology for the current target pair
relevant_micro = df_micro_filtered[
    (df_micro_filtered['ORG_NAME'] == TARGET_ORG) &
    (df_micro_filtered['AB_NAME'] == TARGET_AB_NAME)
].copy()

if relevant_micro.empty:
    print(f"No data found for {pair_label} after initial filtering.␣
↳Skipping.")
    results_store[pair_label] = {'status': 'Skipped - No Micro Data'}
    continue

target_col = 'INTERPRETATION'
relevant_micro = relevant_micro[['SUBJECT_ID', 'HADM_ID', 'CHARTTIME',␣
↳'SPEC_TYPE_DESC', target_col]].dropna(subset=['HADM_ID', 'CHARTTIME',␣
↳target_col]).rename(columns={'CHARTTIME': 'CULTURE_CHARTTIME'})
relevant_micro['HADM_ID'] = relevant_micro['HADM_ID'].astype(int)

# 2. Merge basic patient/admission info
base_df = pd.merge(relevant_micro, df_admissions[['SUBJECT_ID', 'HADM_ID',␣
↳'ADMITTIME', 'ADMISSION_TYPE', 'ETHNICITY']], on=['SUBJECT_ID', 'HADM_ID'],␣
↳how='inner')
base_df = pd.merge(base_df, df_patients[['SUBJECT_ID', 'GENDER', 'DOB']],␣
↳on='SUBJECT_ID', how='inner')

if base_df.empty:
    print(f"No matching admission/patient data found for {pair_label}.␣
↳Skipping.")
    results_store[pair_label] = {'status': 'Skipped - No Adm/Pat Data'}
    del relevant_micro
    gc.collect()
    continue

# --- Age Calculation ---
valid_dates_mask = base_df['ADMITTIME'].notna() & base_df['DOB'].notna()
age_years = (base_df.loc[valid_dates_mask, 'ADMITTIME'].dt.year -
             base_df.loc[valid_dates_mask, 'DOB'].dt.year)
base_df['AGE_AT_ADMISSION'] = np.nan
base_df.loc[valid_dates_mask, 'AGE_AT_ADMISSION'] = age_years
base_df['AGE_AT_ADMISSION'] = base_df['AGE_AT_ADMISSION'].clip(lower=0,␣
↳upper=90)
if base_df['AGE_AT_ADMISSION'].isnull().any():
    median_age = base_df['AGE_AT_ADMISSION'].median()
    if pd.notna(median_age):

```

```

        print(f"Filling {base_df['AGE_AT_ADMISSION'].isnull().sum()} NaN_
↳ages with median age: {median_age:.1f}")
        base_df['AGE_AT_ADMISSION'].fillna(median_age, inplace=True)
    else:
        print(f"Warning: Could not calculate median age for {pair_label}._
↳Filling NaNs with 0.")
        base_df['AGE_AT_ADMISSION'].fillna(0, inplace=True)
    # --- End of Age Calculation ---

    base_df['TIME_WINDOW_END'] = base_df['CULTURE_CHARTTIME'] +_
↳timedelta(hours=TIME_WINDOW_HOURS)
    base_df['PRIOR_ABX_WINDOW_START'] = base_df['CULTURE_CHARTTIME'] -_
↳timedelta(days=PRIOR_ANTIBIOTIC_LOOKBACK_DAYS)

    # 3. Process Labs
    lab_itemids = {'CREATININE': [50912, 50811], 'WBC': [51301], 'BICARBONATE':_
↳[50882]}
    lab_features = []
    current_hadm_ids = base_df['HADM_ID'].unique()
    df_labevents_filtered = df_labevents_full[df_labevents_full['HADM_ID'].
↳isin(current_hadm_ids) & df_labevents_full['VALUENUM'].notna()]
    if not df_labevents_filtered.empty:
        df_labevents_filtered['HADM_ID'] = df_labevents_filtered['HADM_ID'].
↳astype(int)
        merged_labs = pd.merge(base_df[['HADM_ID', 'CULTURE_CHARTTIME',_
↳'TIME_WINDOW_END']], df_labevents_filtered[['HADM_ID', 'ITEMID',_
↳'CHARTTIME', 'VALUENUM']],
                                on='HADM_ID', how='left')
        valid_labs = merged_labs[
            (merged_labs['CHARTTIME'] >= merged_labs['CULTURE_CHARTTIME']) &
            (merged_labs['CHARTTIME'] <= merged_labs['TIME_WINDOW_END']) &
            merged_labs['CHARTTIME'].notna()
        ]
        aggregated_labs = pd.DataFrame(index=pd.Index(current_hadm_ids,_
↳name='HADM_ID'))
        for name, itemids in lab_itemids.items():
            feature_name = f'AVG_{name}_FIRST24H'
            lab_subset = valid_labs[valid_labs['ITEMID'].isin(itemids)]
            mean_vals = lab_subset.groupby('HADM_ID')['VALUENUM'].mean()
            aggregated_labs[feature_name] = mean_vals
            lab_features.append(feature_name)

    base_df = pd.merge(base_df, aggregated_labs, left_on='HADM_ID',_
↳right_index=True, how='left')
    del merged_labs, valid_labs, df_labevents_filtered, aggregated_labs

```

```

else:
    print(f"No relevant lab events found for HADM_IDs in {pair_label}.")
    for name in lab_itemids.keys():
        feature_name = f'AVG_{name}_FIRST24H'
        base_df[feature_name] = np.nan
        lab_features.append(feature_name)

# 4. Process Vitals (REMOVED due to memory constraints)
vital_features = []

# 5. Check for Prior Antibiotics
abx_features = ['HAD_PRIOR_ANTIBIOTICS']
df_prescriptions_filtered = _
df_prescriptions_full[df_prescriptions_full['HADM_ID'].
isin(current_hadm_ids)]
if not df_prescriptions_filtered.empty:
    df_prescriptions_filtered = _
df_prescriptions_filtered[df_prescriptions_filtered['DRUG_TYPE'].str.lower().
isin(['main', 'antibiotic'])]
merged_prescriptions = pd.merge(base_df[['HADM_ID', _
'CULTURE_CHARTTIME', 'PRIOR_ABX_WINDOW_START']].drop_duplicates(),
df_prescriptions_filtered[['HADM_ID', _
'STARTDATE']],
on='HADM_ID', how='left')
valid_prescriptions = merged_prescriptions[
(merged_prescriptions['STARTDATE'] >=_
merged_prescriptions['PRIOR_ABX_WINDOW_START']) &
(merged_prescriptions['STARTDATE'] <=_
merged_prescriptions['CULTURE_CHARTTIME']) &
merged_prescriptions['STARTDATE'].notna()
]
prior_abx_flags = valid_prescriptions[['HADM_ID']].drop_duplicates()
prior_abx_flags['HAD_PRIOR_ANTIBIOTICS'] = 1
base_df = pd.merge(base_df, prior_abx_flags, on='HADM_ID', how='left')
del merged_prescriptions, valid_prescriptions, prior_abx_flags
base_df['HAD_PRIOR_ANTIBIOTICS'] = base_df['HAD_PRIOR_ANTIBIOTICS'].
fillna(0).astype(int)
del df_prescriptions_filtered

# 6. Check for Comorbidities
comorbidity_features = ['HAS_DIABETES']
df_diagnoses_filtered = df_diagnoses_full[df_diagnoses_full['HADM_ID'].
isin(current_hadm_ids) & df_diagnoses_full['ICD9_CODE'].notna()]
if not df_diagnoses_filtered.empty:
    df_diagnoses_filtered['DIABETES_FLAG'] = _
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)

```

```

        diabetes_flags = df_diagnoses_filtered.
↳groupby('HADM_ID')['DIABETES_FLAG'].max().reset_index().
↳rename(columns={'DIABETES_FLAG': 'HAS_DIABETES'})
        base_df = pd.merge(base_df, diabetes_flags, on='HADM_ID', how='left')
        del diabetes_flags
        base_df['HAS_DIABETES'] = base_df['HAS_DIABETES'].fillna(0).astype(int)
        del df_diagnoses_filtered

# 7. Final DataFrame Assembly
required_cols = [
    'SUBJECT_ID', 'HADM_ID', 'CULTURE_CHARTTIME', 'SPEC_TYPE_DESC',
↳target_col,
    'GENDER', 'ADMISSION_TYPE', 'ETHNICITY', 'AGE_AT_ADMISSION'
] + lab_features + vital_features + abx_features + comorbidity_features

for col in required_cols:
    if col not in base_df.columns:
        print(f"Warning: Column {col} not found for pair {pair_label}.
↳Adding as NaN.")
        base_df[col] = np.nan

df_final = base_df[required_cols].copy()
df_final = df_final.drop_duplicates(subset=['HADM_ID', 'CULTURE_CHARTTIME'])
df_final.dropna(subset=[target_col], inplace=True)

print(f"Finished processing for {pair_label}. Final DataFrame shape:
↳{df_final.shape}")

# --- Check for sufficient data and classes BEFORE splitting ---
y_raw_counts = df_final[target_col].value_counts()
if df_final.empty or len(y_raw_counts) < 2:
    print(f"Not enough data or only one class present for {pair_label}.
↳BEFORE splitting. Skipping modeling.")
    results_store[pair_label] = {'status': 'Skipped - Insufficient Data/
↳Classes Pre-Split'}
    del relevant_micro, base_df, df_final
    gc.collect()
    continue
# --- End Check ---

# --- 2. Exploratory Data Analysis (EDA) ---
print(f"\n--- Starting EDA for {pair_label} ---")
print(f"\nTarget variable distribution ({target_col}):")
print(y_raw_counts.to_dict())

```



```

    identifier_cols = ['SUBJECT_ID', 'HADM_ID', 'CULTURE_CHARTTIME']
    all_features = [col for col in df_final.columns if col not in identifier_cols + [target_col]]
    numerical_features = df_final[all_features].select_dtypes(include=np.
    number).columns.tolist()
    categorical_features = df_final[all_features].
    select_dtypes(include=['object', 'category']).columns.tolist()

    # --- Outlier Handling ---
    print(f"\n--- Applying Outlier Capping for {pair_label} ---")
    df_processed = df_final.copy()
    for col in numerical_features:
        if df_processed[col].nunique(dropna=False) <= 2: continue
        if df_processed[col].isnull().all(): continue
        Q1 = df_processed[col].quantile(0.25)
        Q3 = df_processed[col].quantile(0.75)
        IQR = Q3 - Q1
        if pd.isna(IQR) or IQR == 0: continue
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        if pd.isna(lower_bound) or pd.isna(upper_bound): continue
        outliers_low = (df_processed[col] < lower_bound).fillna(False)
        outliers_high = (df_processed[col] > upper_bound).fillna(False)
        num_outliers = outliers_low.sum() + outliers_high.sum()
        if num_outliers > 0:
            df_processed[col] = df_processed[col].mask(outliers_low,
            lower_bound)
            df_processed[col] = df_processed[col].mask(outliers_high,
            upper_bound)

    # --- 3. Preprocessing ---
    print(f"\n--- Starting Preprocessing for {pair_label} ---")
    X = df_processed[all_features]
    y_raw = df_processed[target_col]
    y = target_encoder.transform(y_raw) # Encode target

    # <<< Check minimum class count in the entire dataset y for this pair >>>
    y_counts = Counter(y)
    min_class_count_full = y_counts.most_common()[-1][1] if y_counts else 0

    # Check if minimum count is less than N_SPLITS_CV (needed for
    stratification)
    if min_class_count_full < N_SPLITS_CV:
        print(f"Class with minimum count ({min_class_count_full}) is less than
        N_SPLITS_CV ({N_SPLITS_CV}) for {pair_label} in the full dataset. Skipping
        modeling.")

```

```

        results_store[pair_label] = {'status': f'Skipped - Min Class Count < {N_SPLITS_CV} (Full Set)'}
        # Clean up intermediate dfs specific to this iteration
        del X, y_raw, y, df_processed, df_final, relevant_micro, base_df
        gc.collect()
        continue # Move to the next pair
    # <<< End Check >>>

    numerical_features = X.select_dtypes(include=np.number).columns.tolist()
    categorical_features = X.select_dtypes(include=['object', 'category']).
    columns.tolist()

    numerical_transformer = Pipeline(steps=[('imputer',
    SimpleImputer(strategy='median')), ('scaler', StandardScaler())])
    categorical_transformer = Pipeline(steps=[('imputer',
    SimpleImputer(strategy='most_frequent')), ('onehot',
    OneHotEncoder(handle_unknown='ignore', sparse_output=False))])
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numerical_transformer, [f for f in numerical_features if f
    in X.columns]),
            ('cat', categorical_transformer, [f for f in categorical_features
    if f in X.columns])
        ],
        remainder='passthrough',
        verbose_feature_names_out=False
    )
    preprocessor.set_output(transform='pandas')

    # --- 4. Model Training & Validation ---
    print(f"\n--- Starting Model Training & Validation for {pair_label} ---")
    # Stratify based on the encoded numerical target 'y'
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    random_state=42, stratify=y)

    # Double-check split didn't create single-class sets (less likely now with
    pre-check, but safe)
    if len(np.unique(y_train)) < 2 or len(np.unique(y_test)) < 2:
        print(f"Only one class present in train/test split for {pair_label}
    after encoding. Skipping modeling.")
        results_store[pair_label] = {'status': 'Skipped - Single Class Split'}
        del X, y_raw, y, X_train, X_test, y_train, y_test, df_processed,
    df_final, relevant_micro, base_df
        gc.collect()
        continue

```

```

model_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
↳ ('classifier', RandomForestClassifier(n_estimators=100, random_state=42,
↳ class_weight='balanced', n_jobs=-1))])

print("Performing Cross-Validation...")
cv = StratifiedKFold(n_splits=N_SPLITS_CV, shuffle=True, random_state=42)
mean_cv_f1_weighted = np.nan
# Check minimum class count in y_train BEFORE attempting CV
y_train_counts = Counter(y_train)
min_train_class_count = y_train_counts.most_common()[-1][1] if
↳ y_train_counts else 0

# No need to check len(y_train_counts) > 1 because the earlier check
↳ ensures >= 2 classes
if min_train_class_count >= N_SPLITS_CV:
    try:
        cv_scores = cross_val_score(model_pipeline, X_train, y_train,
↳ cv=cv, scoring='f1_weighted', n_jobs=-1, error_score='raise')
        mean_cv_f1_weighted = np.mean(cv_scores)
        print(f"Mean CV F1-Weighted: {mean_cv_f1_weighted:.4f}")
        cv_status = 'CV Completed'
    except Exception as e:
        print(f"Cross-validation failed for {pair_label} despite check:
↳ {e}")
        cv_status = f'CV Failed: {e}' # Record failure reason
    else:
        print(f"Skipping CV for {pair_label}: Not enough samples in minority
↳ class ({min_train_class_count}) for {N_SPLITS_CV} splits in training data.")
        cv_status = f'Skipped CV - Min Class < {N_SPLITS_CV}' # Record skip
↳ reason

print("Training final model...")
model_pipeline.fit(X_train, y_train)

# --- 5. Evaluation ---
print(f"\n--- Evaluating Model for {pair_label} ---")
y_pred = model_pipeline.predict(X_test)
test_f1_weighted = np.nan

unique_labels_test, unique_counts_test = np.unique(y_test,
↳ return_counts=True)
unique_target_names_test = target_encoder.
↳ inverse_transform(unique_labels_test)

print("\nClassification Report:")

```

```

if len(unique_labels_test) > 1:
    print(classification_report(y_test, y_pred,
                                labels=unique_labels_test,
                                target_names=unique_target_names_test,
                                zero_division=0))
    test_f1_weighted = f1_score(y_test, y_pred, average='weighted',
    ↪ labels=unique_labels_test)
    print(f"Test Set F1-Weighted Score: {test_f1_weighted:.4f}")
else:
    print("Skipping classification report: Only one class in test data.")

# Confusion Matrix for Multiclass
print("\nConfusion Matrix:")
if len(unique_labels_test) > 0: # Ensure there's at least one class to plot
    cm = confusion_matrix(y_test, y_pred, labels=unique_labels_test)
    plt.figure(figsize=(max(6, len(unique_target_names_test)*2), max(4,
    ↪ len(unique_target_names_test)*1.5)))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                 xticklabels=unique_target_names_test,
                 yticklabels=unique_target_names_test)
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.title(f'Confusion Matrix for {pair_label}')
    plt.show()
else:
    print("Skipping confusion matrix: No data in test set.")

# --- 6. Insight Generation ---
print(f"\n--- Generating Insights for {pair_label} ---")
top_features = []
feature_importance_df = None
try:
    feature_names = model_pipeline[:-1].get_feature_names_out()
    importances = model_pipeline.named_steps['classifier'].
    ↪ feature_importances_
    if len(feature_names) == len(importances):
        feature_importance_df = pd.DataFrame({'Feature': feature_names,
        ↪ 'Importance': importances}).sort_values(by='Importance', ascending=False)
        print("\nTop 5 Most Important Features:")
        print(feature_importance_df.head(5))
        top_features = feature_importance_df.head(5)['Feature'].tolist()
    else:
        print(f"Warning: Feature name ({len(feature_names)}) / importance_
        ↪ ({len(importances)}) length mismatch.")
except Exception as e:

```

```

        print(f"Could not extract feature importances: {e}")

    # --- Store Results ---
    # Determine final status based on whether CV was skipped/failed or completed
    final_status = cv_status if cv_status.startswith('Skipped') or cv_status.
↳startswith('CV Failed') else 'Completed'

    results_store[pair_label] = {
        'status': final_status,
        'n_samples': len(df_final),
        'target_distribution': df_final[target_col].
↳value_counts(normalize=True).to_dict(),
        'mean_cv_f1_weighted': mean_cv_f1_weighted, # Will be NaN if CV skipped/
↳failed
        'test_f1_weighted': test_f1_weighted,
        'top_features': top_features,
        'feature_importances': feature_importance_df
    }

    # --- Clean up memory ---
    del X, y_raw, y, X_train, X_test, y_train, y_test, model_pipeline,
↳df_processed, df_final
    del relevant_micro, base_df, feature_importance_df
    gc.collect()

# --- End of Loop ---
print("\n===== Analysis Loop Complete =====")

```

--- Starting Analysis Loop for Top Pairs ---

===== Processing Pair 9621/10: STAPH AUREUS COAG + / OXACILLIN
=====

Processing data for STAPH AUREUS COAG + / OXACILLIN...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

    df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)

```

Finished processing for STAPH AUREUS COAG + / OXACILLIN. Final DataFrame shape:
(7553, 14)

--- Starting EDA for STAPH AUREUS COAG + / OXACILLIN ---

Target variable distribution (INTERPRETATION):

```
{'R': 4499, 'S': 3054, 'P': 0, 'I': 0}
```

--- Applying Outlier Capping for STAPH AUREUS COAG + / OXACILLIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / OXACILLIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / OXACILLIN ---
Performing Cross-Validation...

```
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_diagnoses_filtered['DIABETES_FLAG'] =  
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)
```

Mean CV F1-Weighted: 0.6938

Training final model...

--- Evaluating Model for STAPH AUREUS COAG + / OXACILLIN ---


Classification Report:

	precision	recall	f1-score	support
R	0.74	0.85	0.79	1125
S	0.72	0.55	0.63	764
accuracy			0.73	1889
macro avg	0.73	0.70	0.71	1889
weighted avg	0.73	0.73	0.72	1889

Test Set F1-Weighted Score: 0.7249

Confusion Matrix:

Confusion Matrix for STAPH AUREUS COAG + / OXACILLIN



--- Generating Insights for STAPH AUREUS COAG + / OXACILLIN ---

Top 5 Most Important Features:

	Feature	Importance
0	AGE_AT_ADMISSION	0.199448
2	AVG_WBC_FIRST24H	0.184009
1	AVG_CREATININE_FIRST24H	0.169929
3	AVG_BICARBONATE_FIRST24H	0.168549
5	HAS_DIABETES	0.022445

=====
Processing Pair 9615/10: STAPH AUREUS COAG + / GENTAMICIN
=====

Processing data for STAPH AUREUS COAG + / GENTAMICIN...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_labevents_filtered['HADM_ID'] =  
df_labevents_filtered['HADM_ID'].astype(int)
```

```
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)
```

Finished processing for STAPH AUREUS COAG + / GENTAMICIN. Final DataFrame shape: (6742, 14)

--- Starting EDA for STAPH AUREUS COAG + / GENTAMICIN ---

Target variable distribution (INTERPRETATION):
{ 'S': 6521, 'R': 199, 'I': 22, 'P': 0 }

--- Applying Outlier Capping for STAPH AUREUS COAG + / GENTAMICIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / GENTAMICIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / GENTAMICIN ---

Performing Cross-Validation...

Mean CV F1-Weighted: 0.9666

Training final model...


--- Evaluating Model for STAPH AUREUS COAG + / GENTAMICIN ---

Classification Report:

	precision	recall	f1-score	support
I	0.00	0.00	0.00	5
R	0.93	0.28	0.43	50
S	0.98	1.00	0.99	1631
accuracy			0.97	1686
macro avg	0.64	0.43	0.47	1686
weighted avg	0.97	0.97	0.97	1686

Test Set F1-Weighted Score: 0.9676

Confusion Matrix:



--- Generating Insights for STAPH AUREUS COAG + / GENTAMICIN ---

Top 5 Most Important Features:

	Feature	Importance
1	AVG_CREATININE_FIRST24H	0.167109
0	AGE_AT_ADMISSION	0.165230
2	AVG_WBC_FIRST24H	0.151905
3	AVG_BICARBONATE_FIRST24H	0.135805
4	HAD_PRIOR_ANTIBIOTICS	0.047818

=====
Processing Pair 9617/10: STAPH AUREUS COAG + / LEVOFLOXACIN
=====

Processing data for STAPH AUREUS COAG + / LEVOFLOXACIN...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)

Finished processing for STAPH AUREUS COAG + / LEVOFLOXACIN. Final DataFrame
shape: (6717, 14)

--- Starting EDA for STAPH AUREUS COAG + / LEVOFLOXACIN ---

Target variable distribution (INTERPRETATION):
{'R': 4158, 'S': 2463, 'I': 96, 'P': 0}

--- Applying Outlier Capping for STAPH AUREUS COAG + / LEVOFLOXACIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / LEVOFLOXACIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / LEVOFLOXACIN
---
Performing Cross-Validation...
Mean CV F1-Weighted: 0.7216
Training final model...

--- Evaluating Model for STAPH AUREUS COAG + / LEVOFLOXACIN ---

Classification Report:

```

	precision	recall	f1-score	support
I	0.50	0.17	0.25	24
R	0.74	0.89	0.81	1040
S	0.72	0.50	0.59	616
accuracy			0.73	1680
macro avg	0.65	0.52	0.55	1680
weighted avg	0.73	0.73	0.72	1680


```

Test Set F1-Weighted Score: 0.7177

Confusion Matrix:

```

Confusion Matrix for STAPH AUREUS COAG + / LEVOFLOXACIN



--- Generating Insights for STAPH AUREUS COAG + / LEVOFLOXACIN ---

Top 5 Most Important Features:

	Feature	Importance
2	AVG_WBC_FIRST24H	0.182624
0	AGE_AT_ADMISSION	0.179135
1	AVG_CREATININE_FIRST24H	0.164688
3	AVG_BICARBONATE_FIRST24H	0.156042
4	HAD_PRIOR_ANTIBIOTICS	0.031754

=====
Processing Pair 9614/10: STAPH AUREUS COAG + / ERYTHROMYCIN
=====

Processing data for STAPH AUREUS COAG + / ERYTHROMYCIN...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)

Finished processing for STAPH AUREUS COAG + / ERYTHROMYCIN. Final DataFrame
shape: (6479, 14)

--- Starting EDA for STAPH AUREUS COAG + / ERYTHROMYCIN ---

Target variable distribution (INTERPRETATION):
{'R': 4515, 'S': 1829, 'I': 135, 'P': 0}

--- Applying Outlier Capping for STAPH AUREUS COAG + / ERYTHROMYCIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / ERYTHROMYCIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / ERYTHROMYCIN
---
Performing Cross-Validation...
Mean CV F1-Weighted: 0.6915
Training final model...

--- Evaluating Model for STAPH AUREUS COAG + / ERYTHROMYCIN ---

Classification Report:

```


	precision	recall	f1-score	support
I	0.47	0.21	0.29	34
R	0.77	0.93	0.84	1129
S	0.69	0.36	0.47	457
accuracy			0.75	1620
macro avg	0.64	0.50	0.53	1620
weighted avg	0.74	0.75	0.73	1620

```

Test Set F1-Weighted Score: 0.7262

Confusion Matrix:

```



--- Generating Insights for STAPH AUREUS COAG + / ERYTHROMYCIN ---

Top 5 Most Important Features:

	Feature	Importance
2	AVG_WBC_FIRST24H	0.185675
0	AGE_AT_ADMISSION	0.178579
3	AVG_BICARBONATE_FIRST24H	0.170249
1	AVG_CREATININE_FIRST24H	0.165571
5	HAS_DIABETES	0.026010

=====
Processing Pair 9612/10: STAPH AUREUS COAG + / CLINDAMYCIN
=====

Processing data for STAPH AUREUS COAG + / CLINDAMYCIN...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)

Finished processing for STAPH AUREUS COAG + / CLINDAMYCIN. Final DataFrame
shape: (5028, 14)

--- Starting EDA for STAPH AUREUS COAG + / CLINDAMYCIN ---

Target variable distribution (INTERPRETATION):
{'R': 2679, 'S': 2347, 'I': 2, 'P': 0}

--- Applying Outlier Capping for STAPH AUREUS COAG + / CLINDAMYCIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / CLINDAMYCIN ---
Class with minimum count (2) is less than N_SPLITS_CV (5) for STAPH AUREUS COAG
+ / CLINDAMYCIN in the full dataset. Skipping modeling.

===== Processing Pair 9627/10: STAPH AUREUS COAG + / TETRACYCLINE
=====

Processing data for STAPH AUREUS COAG + / TETRACYCLINE...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)

Finished processing for STAPH AUREUS COAG + / TETRACYCLINE. Final DataFrame
shape: (4193, 14)

```

```

--- Starting EDA for STAPH AUREUS COAG + / TETRACYCLINE ---

Target variable distribution (INTERPRETATION):
{'S': 3972, 'R': 220, 'I': 1, 'P': 0}

--- Applying Outlier Capping for STAPH AUREUS COAG + / TETRACYCLINE ---

--- Starting Preprocessing for STAPH AUREUS COAG + / TETRACYCLINE ---
Class with minimum count (1) is less than N_SPLITS_CV (5) for STAPH AUREUS COAG
+ / TETRACYCLINE in the full dataset. Skipping modeling.

===== Processing Pair 9622/10: STAPH AUREUS COAG + / PENICILLIN
=====

Processing data for STAPH AUREUS COAG + / PENICILLIN...

/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)

Finished processing for STAPH AUREUS COAG + / PENICILLIN. Final DataFrame shape:
(4069, 14)

--- Starting EDA for STAPH AUREUS COAG + / PENICILLIN ---

Target variable distribution (INTERPRETATION):
{'R': 4003, 'S': 66, 'P': 0, 'I': 0}

--- Applying Outlier Capping for STAPH AUREUS COAG + / PENICILLIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / PENICILLIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / PENICILLIN
---
Performing Cross-Validation...

```

Mean CV F1-Weighted: 0.9785

Training final model...


--- Evaluating Model for STAPH AUREUS COAG + / PENICILLIN ---

Classification Report:

	precision	recall	f1-score	support
R	0.99	1.00	0.99	1001
S	1.00	0.24	0.38	17
accuracy			0.99	1018
macro avg	0.99	0.62	0.69	1018
weighted avg	0.99	0.99	0.98	1018

Test Set F1-Weighted Score: 0.9833

Confusion Matrix:



--- Generating Insights for STAPH AUREUS COAG + / PENICILLIN ---

Top 5 Most Important Features:

Feature	Importance
---------	------------

2	AVG_WBC_FIRST24H	0.162217
0	AGE_AT_ADMISSION	0.156591
3	AVG_BICARBONATE_FIRST24H	0.151228
1	AVG_CREATININE_FIRST24H	0.140259
4	HAD_PRIOR_ANTIBIOTICS	0.079369

```
===== Processing Pair 9630/10: STAPH AUREUS COAG + / VANCOMYCIN
=====
```

Processing data for STAPH AUREUS COAG + / VANCOMYCIN...

```
/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)
```

Finished processing for STAPH AUREUS COAG + / VANCOMYCIN. Final DataFrame shape: (4123, 14)

--- Starting EDA for STAPH AUREUS COAG + / VANCOMYCIN ---

```
Target variable distribution (INTERPRETATION):
{'S': 4113, 'I': 10, 'P': 0, 'R': 0}
```

--- Applying Outlier Capping for STAPH AUREUS COAG + / VANCOMYCIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / VANCOMYCIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / VANCOMYCIN ---

```
Performing Cross-Validation...
Mean CV F1-Weighted: 0.9970
Training final model...
```


--- Evaluating Model for STAPH AUREUS COAG + / VANCOMYCIN ---

Classification Report:

	precision	recall	f1-score	support
I	0.40	0.67	0.50	3
S	1.00	1.00	1.00	1028
accuracy			1.00	1031
macro avg	0.70	0.83	0.75	1031
weighted avg	1.00	1.00	1.00	1031

Test Set F1-Weighted Score: 0.9966

Confusion Matrix:



--- Generating Insights for STAPH AUREUS COAG + / VANCOMYCIN ---

Top 5 Most Important Features:

	Feature	Importance
1	AVG_CREATININE_FIRST24H	0.228347
3	AVG_BICARBONATE_FIRST24H	0.142716
2	AVG_WBC_FIRST24H	0.136160
0	AGE_AT_ADMISSION	0.125226
10	SPEC_TYPE_DESC_BLOOD CULTURE	0.088965

```
===== Processing Pair 9626/10: STAPH AUREUS COAG + / RIFAMPIN
=====
```

Processing data for STAPH AUREUS COAG + / RIFAMPIN..

```
/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
```

```
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)
```

Finished processing for STAPH AUREUS COAG + / RIFAMPIN. Final DataFrame shape: (3962, 14)

--- Starting EDA for STAPH AUREUS COAG + / RIFAMPIN ---

Target variable distribution (INTERPRETATION):
{ 'S': 3850, 'R': 86, 'I': 26, 'P': 0 }

--- Applying Outlier Capping for STAPH AUREUS COAG + / RIFAMPIN ---

--- Starting Preprocessing for STAPH AUREUS COAG + / RIFAMPIN ---

--- Starting Model Training & Validation for STAPH AUREUS COAG + / RIFAMPIN ---

Performing Cross-Validation...

Mean CV F1-Weighted: 0.9647

Training final model...

--- Evaluating Model for STAPH AUREUS COAG + / RIFAMPIN ---

Classification Report:


	precision	recall	f1-score	support
I	1.00	0.17	0.29	6
R	0.83	0.45	0.59	22
S	0.98	1.00	0.99	963
accuracy			0.98	991

macro avg	0.94	0.54	0.62	991
weighted avg	0.98	0.98	0.98	991

Test Set F1-Weighted Score: 0.9770

Confusion Matrix:

Confusion Matrix for STAPH AUREUS COAG + / RIFAMPIN



--- Generating Insights for STAPH AUREUS COAG + / RIFAMPIN ---

Top 5 Most Important Features:

	Feature	Importance
2	AVG_WBC_FIRST24H	0.167343
0	AGE_AT_ADMISSION	0.163112
3	AVG_BICARBONATE_FIRST24H	0.162699
1	AVG_CREATININE_FIRST24H	0.139951
5	HAS_DIABETES	0.040820

=====
 Processing Pair 4065/10: ESCHERICHIA COLI / GENTAMICIN
 =====

Processing data for ESCHERICHIA COLI / GENTAMICIN...

```
/var/tmp/ipykernel_16006/2980027711.py:64: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_labevents_filtered['HADM_ID'] =
df_labevents_filtered['HADM_ID'].astype(int)
/var/tmp/ipykernel_16006/2980027711.py:117: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_diagnoses_filtered['DIABETES_FLAG'] =
df_diagnoses_filtered['ICD9_CODE'].str.startswith('250').astype(int)
```

Finished processing for ESCHERICHIA COLI / GENTAMICIN. Final DataFrame shape: (3505, 14)

--- Starting EDA for ESCHERICHIA COLI / GENTAMICIN ---

Target variable distribution (INTERPRETATION):
{ 'S': 3045, 'R': 429, 'I': 31, 'P': 0 }

--- Applying Outlier Capping for ESCHERICHIA COLI / GENTAMICIN ---

--- Starting Preprocessing for ESCHERICHIA COLI / GENTAMICIN ---

--- Starting Model Training & Validation for ESCHERICHIA COLI / GENTAMICIN ---
Performing Cross-Validation...
Mean CV F1-Weighted: 0.8565
Training final model...

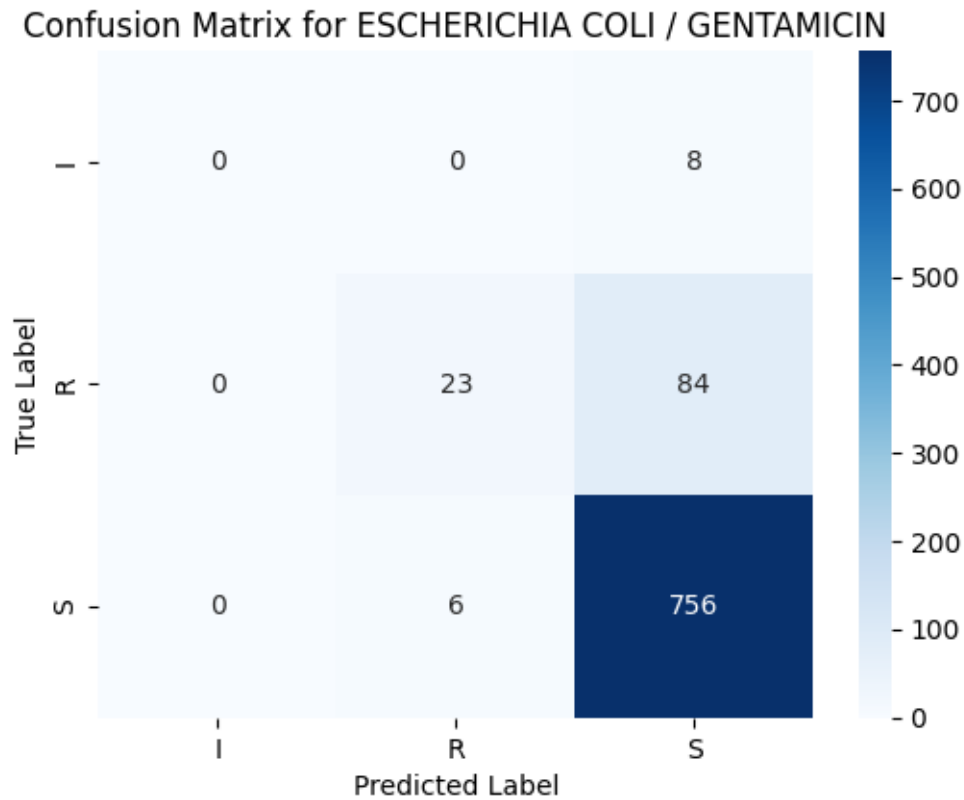
--- Evaluating Model for ESCHERICHIA COLI / GENTAMICIN ---

Classification Report:

	precision	recall	f1-score	support
I	0.00	0.00	0.00	8
R	0.79	0.21	0.34	107
S	0.89	0.99	0.94	762
accuracy			0.89	877
macro avg	0.56	0.40	0.43	877
weighted avg	0.87	0.89	0.86	877

Test Set F1-Weighted Score: 0.8573

Confusion Matrix:



--- Generating Insights for ESCHERICHIA COLI / GENTAMICIN ---

Top 5 Most Important Features:

	Feature	Importance
2	AVG_WBC_FIRST24H	0.169332
0	AGE_AT_ADMISSION	0.156132
3	AVG_BICARBONATE_FIRST24H	0.154887
1	AVG_CREATININE_FIRST24H	0.148444
9	SPEC_TYPE_DESC_BLOOD CULTURE	0.034061

===== Analysis Loop Complete =====

```
[53]: # --- Summarize Results ---
print("\n--- Summary of Results Across Top Pairs ---")
# Create DataFrame from the results store
results_summary_df = pd.DataFrame.from_dict(results_store, orient='index')

# Define expected columns and ensure they exist, adding NA if missing
```

```

expected_cols = ['status', 'n_samples', 'target_distribution',
↳ 'mean_cv_f1_weighted', 'test_f1_weighted', 'top_features']
for col in expected_cols:
    if col not in results_summary_df.columns:
        results_summary_df[col] = pd.NA

# Select and reorder columns
results_summary_df = results_summary_df[expected_cols]

# --- Enhanced Formatting and Explanation ---
# Create a copy for display formatting
results_display_df = results_summary_df.copy()

# Format target distribution for readability
def format_distribution(dist_dict):
    if pd.isna(dist_dict) or not isinstance(dist_dict, dict):
        return "N/A"
    return ', '.join([f"{k}: {v:.1%}" for k, v in sorted(dist_dict.items())])

results_display_df['target_distribution'] =
↳ results_display_df['target_distribution'].apply(format_distribution)

# Format F1 scores
results_display_df['mean_cv_f1_weighted'] =
↳ results_display_df['mean_cv_f1_weighted'].map('{:.4f}'.format).
↳ replace('nan', 'N/A')
results_display_df['test_f1_weighted'] = results_display_df['test_f1_weighted'].
↳ map('{:.4f}'.format).replace('nan', 'N/A')

# Format top features list
def format_top_features(feature_list):
    # Check if input is a list and not empty
    if isinstance(feature_list, list) and feature_list:
        # Truncate long feature names
        truncated_list = [f[:30] + '...' if len(f) > 33 else f for f in
↳ feature_list]
        return ', '.join(truncated_list)
    # Handle cases where feature_list might be NaN, None, or an empty list
    elif pd.isna(feature_list) or not feature_list:
        return "N/A"
    # Fallback for unexpected types (though less likely now)
    else:
        return "Invalid Format"

results_display_df['top_features'] = results_display_df['top_features'].
↳ apply(format_top_features)

```

```

# Rename columns for better display
results_display_df = results_display_df.rename(columns={
    'status': 'Status',
    'n_samples': 'N Samples',
    'target_distribution': 'Target Distr.',
    'mean_cv_f1_weighted': 'Mean CV F1 (W)',
    'test_f1_weighted': 'Test F1 (W)',
    'top_features': 'Top 5 Features (Approx.)'
})

# Sort by test F1 score (descending), putting NaNs/NAs last
# Convert F1 score back to numeric for sorting, handling 'N/A'
results_display_df['Test F1 (W)_sort'] = pd.to_numeric(results_display_df['Test_
    ↪F1 (W)'], errors='coerce')
results_display_df = results_display_df.sort_values(by='Test F1 (W)_sort',
    ↪ascending=False, na_position='last').drop(columns=['Test F1 (W)_sort'])

print("\n--- Formatted Summary ---")
print("Explanation of Columns:")
print("- Status: Outcome of the modeling process for the pair (Completed,
    ↪Skipped, CV Failed, etc.).")
print("- N Samples: Number of valid microbiology events used for modeling this_
    ↪pair.")
print("- Target Distr.: Distribution of target classes (R, S, I) in the dataset_
    ↪for this pair.")
print("- Mean CV F1 (W): Average Weighted F1-score from 5-fold cross-validation_
    ↪(if performed). 'N/A' if skipped or failed.")
print("- Test F1 (W): Weighted F1-score on the held-out test set. 'N/A' if_
    ↪evaluation failed.")
print("- Top 5 Features (Approx.): The most important features identified by_
    ↪the RandomForest model (names might be truncated).")
print("-" * 80) # Separator

# Display the formatted DataFrame
# Adjust display options if needed for wide tables
pd.set_option('display.max_colwidth', None) # Show full content of columns like_
    ↪features
pd.set_option('display.width', 120) # Adjust display width
print(results_display_df)
pd.reset_option('display.max_colwidth')
pd.reset_option('display.width')
print("\n--- Full Analysis Complete ---")

```

--- Summary of Results Across Top Pairs ---

--- Formatted Summary ---

Explanation of Columns:

- Status: Outcome of the modeling process for the pair (Completed, Skipped, CV Failed, etc.).
- N Samples: Number of valid microbiology events used for modeling this pair.
- Target Distr.: Distribution of target classes (R, S, I) in the dataset for this pair.
- Mean CV F1 (W): Average Weighted F1-score from 5-fold cross-validation (if performed). 'N/A' if skipped or failed.
- Test F1 (W): Weighted F1-score on the held-out test set. 'N/A' if evaluation failed.
- Top 5 Features (Approx.): The most important features identified by the RandomForest model (names might be truncated).

		Status	N
Samples \			
STAPH AUREUS COAG + / VANCOMYCIN		Completed	
4123.0			
STAPH AUREUS COAG + / PENICILLIN		Completed	
4069.0			
STAPH AUREUS COAG + / RIFAMPIN		Completed	
3962.0			
STAPH AUREUS COAG + / GENTAMICIN		Completed	
6742.0			
ESCHERICHIA COLI / GENTAMICIN		Completed	
3505.0			
STAPH AUREUS COAG + / ERYTHROMYCIN		Completed	
6479.0			
STAPH AUREUS COAG + / OXACILLIN		Completed	
7553.0			
STAPH AUREUS COAG + / LEVOFLOXACIN		Completed	
6717.0			
STAPH AUREUS COAG + / CLINDAMYCIN	Skipped - Min Class Count < 5 (Full Set)		
NaN			
STAPH AUREUS COAG + / TETRACYCLINE	Skipped - Min Class Count < 5 (Full Set)		
NaN			
		Target Distr.	Mean CV
F1 (W) Test F1 (W) \			
STAPH AUREUS COAG + / VANCOMYCIN	I: 0.2%, P: 0.0%, R: 0.0%, S: 99.8%		
0.9970 0.9966			
STAPH AUREUS COAG + / PENICILLIN	I: 0.0%, P: 0.0%, R: 98.4%, S: 1.6%		
0.9785 0.9833			
STAPH AUREUS COAG + / RIFAMPIN	I: 0.7%, P: 0.0%, R: 2.2%, S: 97.2%		
0.9647 0.9770			
STAPH AUREUS COAG + / GENTAMICIN	I: 0.3%, P: 0.0%, R: 3.0%, S: 96.7%		
0.9666 0.9676			

ESCHERICHIA COLI / GENTAMICIN	I: 0.9%, P: 0.0%, R: 12.2%, S: 86.9%
0.8565 0.8573	
STAPH AUREUS COAG + / ERYTHROMYCIN	I: 2.1%, P: 0.0%, R: 69.7%, S: 28.2%
0.6915 0.7262	
STAPH AUREUS COAG + / OXACILLIN	I: 0.0%, P: 0.0%, R: 59.6%, S: 40.4%
0.6938 0.7249	
STAPH AUREUS COAG + / LEVOFLOXACIN	I: 1.4%, P: 0.0%, R: 61.9%, S: 36.7%
0.7216 0.7177	
STAPH AUREUS COAG + / CLINDAMYCIN	N/A
N/A N/A	
STAPH AUREUS COAG + / TETRACYCLINE	N/A
N/A N/A	

Top 5 Features (Approx.)

STAPH AUREUS COAG + / VANCOMYCIN	AVG_CREATININE_FIRST24H,
AVG_BICARBONATE_FIRST24H, AVG_WBC_FIRST24H, AGE_AT_ADMISSION,	
SPEC_TYPE_DESC_BLOOD CULTURE	
STAPH AUREUS COAG + / PENICILLIN	AVG_WBC_FIRST24H, AGE_AT_ADMISSION,
AVG_BICARBONATE_FIRST24H, AVG_CREATININE_FIRST24H, HAD_PRIOR_ANTIBIOTICS	
STAPH AUREUS COAG + / RIFAMPIN	AVG_WBC_FIRST24H,
AGE_AT_ADMISSION, AVG_BICARBONATE_FIRST24H, AVG_CREATININE_FIRST24H,	
HAS_DIABETES	
STAPH AUREUS COAG + / GENTAMICIN	AVG_CREATININE_FIRST24H,
AGE_AT_ADMISSION, AVG_WBC_FIRST24H, AVG_BICARBONATE_FIRST24H,	
HAD_PRIOR_ANTIBIOTICS	
ESCHERICHIA COLI / GENTAMICIN	AVG_WBC_FIRST24H, AGE_AT_ADMISSION,
AVG_BICARBONATE_FIRST24H, AVG_CREATININE_FIRST24H, SPEC_TYPE_DESC_BLOOD CULTURE	
STAPH AUREUS COAG + / ERYTHROMYCIN	AVG_WBC_FIRST24H,
AGE_AT_ADMISSION, AVG_BICARBONATE_FIRST24H, AVG_CREATININE_FIRST24H,	
HAS_DIABETES	
STAPH AUREUS COAG + / OXACILLIN	AGE_AT_ADMISSION,
AVG_WBC_FIRST24H, AVG_CREATININE_FIRST24H, AVG_BICARBONATE_FIRST24H,	
HAS_DIABETES	
STAPH AUREUS COAG + / LEVOFLOXACIN	AVG_WBC_FIRST24H, AGE_AT_ADMISSION,
AVG_CREATININE_FIRST24H, AVG_BICARBONATE_FIRST24H, HAD_PRIOR_ANTIBIOTICS	
STAPH AUREUS COAG + / CLINDAMYCIN	
N/A	
STAPH AUREUS COAG + / TETRACYCLINE	
N/A	

--- Full Analysis Complete ---

[]:

[]: